

Cloud Technology

Assignment 1

Student Name: Junli Liang

Email: junli.liang4@mail.dcu.ie

Student Number: 18212690

Github link: https://github.com/diaopk/useful_files

Introduction

This report illustrates steps on how to attempt to do the tasks. The structure of this report is divided by 4 sections corresponding to 4 individual tasks. Commands and description will be associated with each task in terms of ways to doing.

Task 1 – Acquire Post Data

Required to query 200,000 posts by ViewCount. 4 queries used to determine 4 * 50,000 posts intervals. The queried data in zip format is available

https://github.com/diaopk/useful_files/blob/master/posts_v2.zip

```
1 /* lower limit - 200,010 posts */
2 select count(*) from posts where posts.ViewCount > 28465;
3
4 /* 1st interval 49,997 posts */
5 select count(*) from posts where posts.viewCount > 28464 and posts.ViewCount < 36444;
6
7 /* 2nd interval 50,001 posts */
8 select count(*) from posts where posts.viewCount > 36444 and posts.ViewCount < 50821;
9
10 /* 3rd interval 50,000 posts */
11 select count(*) from posts where posts.ViewCount > 50821 and posts.ViewCount < 86334;
12
13 /* 4th interval 49,997 posts */
14 select count(*) from posts where posts.ViewCount > 86334;
```

The numbers used to filter the ViewCount are chosen by multiple trials.

Task 2 – Use with Pig to Load/Extract Data

With Pig, I am able to load the raw data and extract it depending on the actual needs. The entire query, pig, hive commands are available at

https://github.com/diaopk/useful_files/blob/master/sql_pig_hive_script.sql. With the following Pig command, queried data can be loaded into Pig. It defines a schema for loading data and specifies “,” format.

```
18 posts1 = load './posts_1.csv' using PigStorage(',') as (id:int, PostTypeId:int, Accept
edAnswerId:int, ParentId:int, CreationDate:datetime, DeletionDate:datetime, Score:int,
ViewCount:int, OwnerUserId:int, OwnerDisplayName:chararray, LastEditorUserId:int, Las
tEditorDisplayName:chararray, LastEditDate:datetime, LastActivityDate:datetime, Title:
chararray, Tags:chararray, AnswerCount:int, CommentCount:int, FavoriteCount:int, Close
dDate:datetime, CommunityOwnedDate:datetime);
```

As there are 4 different csv files I needed to load in so the command above would be run 4 times. `posts = UNION posts1, posts2, posts3, posts4;` command is used to combine 4 loaded data. `FILTER posts BY (Body MATCHES '.*hadoop.*');` command is used to filter posts with the word “Hadoop” is included in Body column. This command is useful for task 3.3 described in the assignment description. Eventually the

result can be stored as files using the command below.

```
STORE filter_posts INTO '<pig_output>' USING PigStorage(',');
```

Task 3 – Use with Hive to Query and Query Result Screenshots

This section uses the files generated by Pig in the previous section in Hive. First of all, I need to create a table for posts/filter_posts. In total there are 22 columns of data.

```
32 CREATE TABLE IF NOT EXISTS posts
33 (id INT, PostTypeId INT, AcceptedAnswerId INT, ParentId INT, CreationDate TIMESTAMP, DeletionDate TIMESTAMP, Score INT, ViewCount INT, Body STRING, OwnerUserId INT, OwnerDisplayName STRING, LastEditorUserId INT, LastEditorDisplayName STRING, LastEditDate TIMESTAMP, LastActivityDate TIMESTAMP, Title STRING, Tags STRING, AnswerCount INT, CommentCount INT, FavoriteCount INT, ClosedDate TIMESTAMP, CommunityOwnedDate TIMESTAMP)
34 COMMENT 'Posts of Stackover Exchange' ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE [LOCATION <input-path>];
```

- Task 3.1 Query the top 10 posts by score.

Query:

```
SELECT id,score FROM posts ORDER BY score DESC LIMIT 10;
```

Result shown PostId and Score:

```
OK
11227809      22586
927358      19063
2003505      14715
292357      10726
477816      9517
231767      8943
1642028      8112
348170      7894
503093      7733
179123      7676
Time taken: 7.516 seconds, Fetched: 10 row(s)
```

- Task 3.2 Query the top 10 users by post score.

Query:

```
SELECT OwnerUserId AS user_id, SUM(score) AS scores FROM posts
WHERE OwnerUserId IS NOT NULL GROUP BY OwnerUserId ORDER BY scores DESC
LIMIT 10;
```

Result shown OwnerUserId and combined Score:

```
OK
87234      32477
4883      22784
9951      22628
6068      21507
89904      19731
51816      16542
49153      15644
95592      15409
63051      14858
39677      14794
Time taken: 6.405 seconds, Fetched: 10 row(s)
```

- Task 3.3 Query the number of distinct users, who used the word “Hadoop” in one of their posts. As the data already is filtered by Pig, then that is relatively easy task with Hive.

Query:

```
SELECT COUNT(DISTINCT OwnerUserId) FROM filter_posts WHERE OwnerUserId IS NOT NULL;
```

Result:

```
OK
88
Time taken: 1.809 seconds, Fetched: 1 row(s)
```

- **Task 3.4 Pre-task of Task 4** – generate a csv format file for MapReduce program. With the command below, I can filter out top 10 users ordered by score and associated Body content. This query concludes subquery, which basically the query result from task 3.2. Then Extract top 10 OwnerUserId order by score and their Body content to generate file.

```
INSERT OVERWRITE LOCAL DIRECTORY <paht-to-output>
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
SELECT owneruserid, body FROM full_posts WHERE owneruserid IN
(SELECT top10user.user_id FROM
  (SELECT owneruserid AS user_id, SUM(score) AS score FROM posts
   WHERE owneruserid IS NOT NULL GROUP BY owneruserid ORDER BY score DESC limit 10)
 AS top10user);
```

Task 4 – Use with MapReduce

The python MapReduce script is available at

https://github.com/diaopk/useful_files/blob/master/mr_count.py. The idea behind is that I used two python datastructure Dictionary to record the TF values and IDF values respectively.

tf_dict is defined in { user_id: { word_a: TF_value,
word_b: TF_value,
... }
}

idf_dict is defined in

```
{ user_id: [ total_number_of_docs, { word_a: [num_doc_with_word_a, idf_value],
word_b: [num_doc_with_word_b, idf_value],
... }
]
}
```

Mapper contains three components, mapper_init(), mapper_middle() and mapper_final(). Mapper_init() initialises two data structures. Mapper_middle() concurrently share the two defined dictionaries and update values such as total number of documents, number of documents with 'word_a/b/...' in it, etc. Mapper_final() passes user id plus a word, separated by '–', and a calculated TF-IDF value as tuple (e.g. (123456-how, 0.xxxx))to combiner(). There is only one step MapReduce job, which may not fully utilise the power of MapReduce. The partial result shown below is user id (51816-<word>, tf_idf_value).

```
"51816-as" 0.0243797137225529
"51816-asked" 0.014479239745963811
"51816-asking" 0.022596389300519285
"51816-assembly" 0.02330252646616051
"51816-assigned" 0.014410208340437844
"51816-assignment" 0.031381812627582654
"51816-at" 0.044504702365528906
"51816-attr" 0.03765817515309918
"51816-attribute" 0.05021090020413224
"51816-attributes" 0.028399506474937497
"51816-author" 0.013080568111761445
"51816-available" 0.031381812627582654
"51816-avoided" 0.02255510870677228
"51816-backwards" 0.02593837501278812
"51816-base" 0.06540492400288636
"51816-based" 0.010277916629025
"51816-basic" 0.012532451208691366
"51816-be" 0.041283938064035464
"51816-because" 0.01803532876015593
"51816-before" 0.012306593187287313
"51816-begin" 0.02330252646616051
"51816-behave" 0.01945378125959109
```

Limitation

As shown, the MapReduce script shows all users and their word TF-IDF values. Results are not displayed in order and not able to show top 10 word-tfidf values, yet. Additional works need to be done in combiner and reducer.